

Databases 1

Daniel POP

Course Outline

- ~~1. Introduction to database approach~~
- ~~2. The database environment~~
- ~~3. Introduction to The Relational Model~~
4. Views
5. Transactions
6. SQL Constraints
- ~~7. Relational Database Design. Theory and practice~~
8. An Introduction to Database Performance. Indexing
9. JSON Support in Relational Database Management Systems
10. NoSQL Databases

A solid blue vertical bar is located on the left side of the slide, extending from the top to the bottom.

Week 10

SQL Constraints

Views

SQL Constraints

- Define “rules” for data in the relational database
 - ensures data consistency and integrity
 - any data action that violates the constraint is rejected
- Classification of constraints
 - Column level
 - Table level
- Common SQL constraints are
 - NOT NULL
 - PRIMARY KEY
 - FOREIGN KEY
 - DEFAULT
 - UNIQUE
 - CHECK
 - INDEX
- Can be defined either at table creation time

```
CREATE TABLE Students (  
    Id INT PRIMARY KEY,  
    Name VARCHAR(100) NOT NULL,  
    Address VARCHAR(256),  
    MajorId INT NOT NULL  
        FOREIGN KEY REFERENCES Major(Id),  
    Active CHAR(1))
```
- Or later

```
ALTER TABLE Students  
    MODIFY Active DEFAULT 'Y';
```

Good coding practice: always give meaningful names to constraints

UNIQUE Constraint

- Ensures that all values in one (or more) column(s) are unique
- Useful to represent other candidate keys of the domain, when surrogate/auto-generated primary keys are used;
- These are the alternate keys, hence it is recommended that their starts with AK_
- While table may have only one PK, multiple unique constraints (alternate candidate keys) can be added
- Example:

```
ALTER TABLE Students
```

```
ADD CONSTRAINT AK_Students_Name_Address UNIQUE (Name, Address);
```

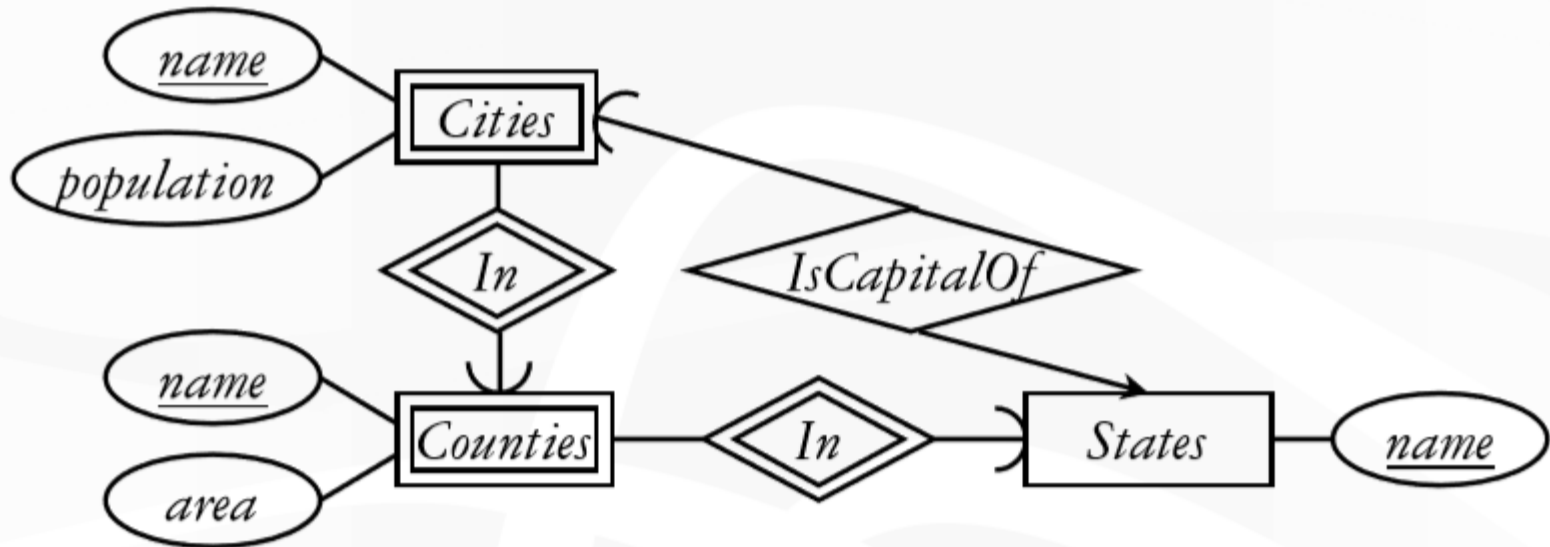
CHECK Constraint

- Limit the values for a specified column or limit the values in certain columns based on other values in other columns of the same row
- The predicate cannot contain queries, but may call user-defined/system functions
- Example:

```
ALTER TABLE Students
```

```
ADD CONSTRAINT CK_Students_Valid_Id_Name CHECK (Id >= 0 AND Name <> '');
```

Example



- ❖ Technically, nothing in this design could prevent a city in state *X* from being the capital of another state *Y*, but oh well...

For demo purpose, assume the following structure for table CapitalOf
CapitalOf(CityId, StateName)
where CityId is a Foreign Key referencing Cities(Id)

Example

```
CREATE FUNCTION IsCityInState
(
    @CityId INT,
    @StateName VARCHAR(128)
)
RETURNS INT
AS
BEGIN
    IF @StateName = (SELECT StateName
                     FROM Counties CN
                     INNER JOIN Cities CT ON CN.Id = CT.CountyId
                     WHERE CT.Id = @CityId)
        return 1
    return 0
END

ALTER TABLE CapitalOf
    WITH CHECK ADD CONSTRAINT CK_ValidState
    CHECK (CityId IS NULL OR IsCityInState(CityId, StateName) = 1)
```


Week 10

Views

Scenario

The mobile dev team of our university develops a student-centred mobile application aiming at helping students to manage their enrollments. How should the app request the list of all enrollments for the logged-in student?

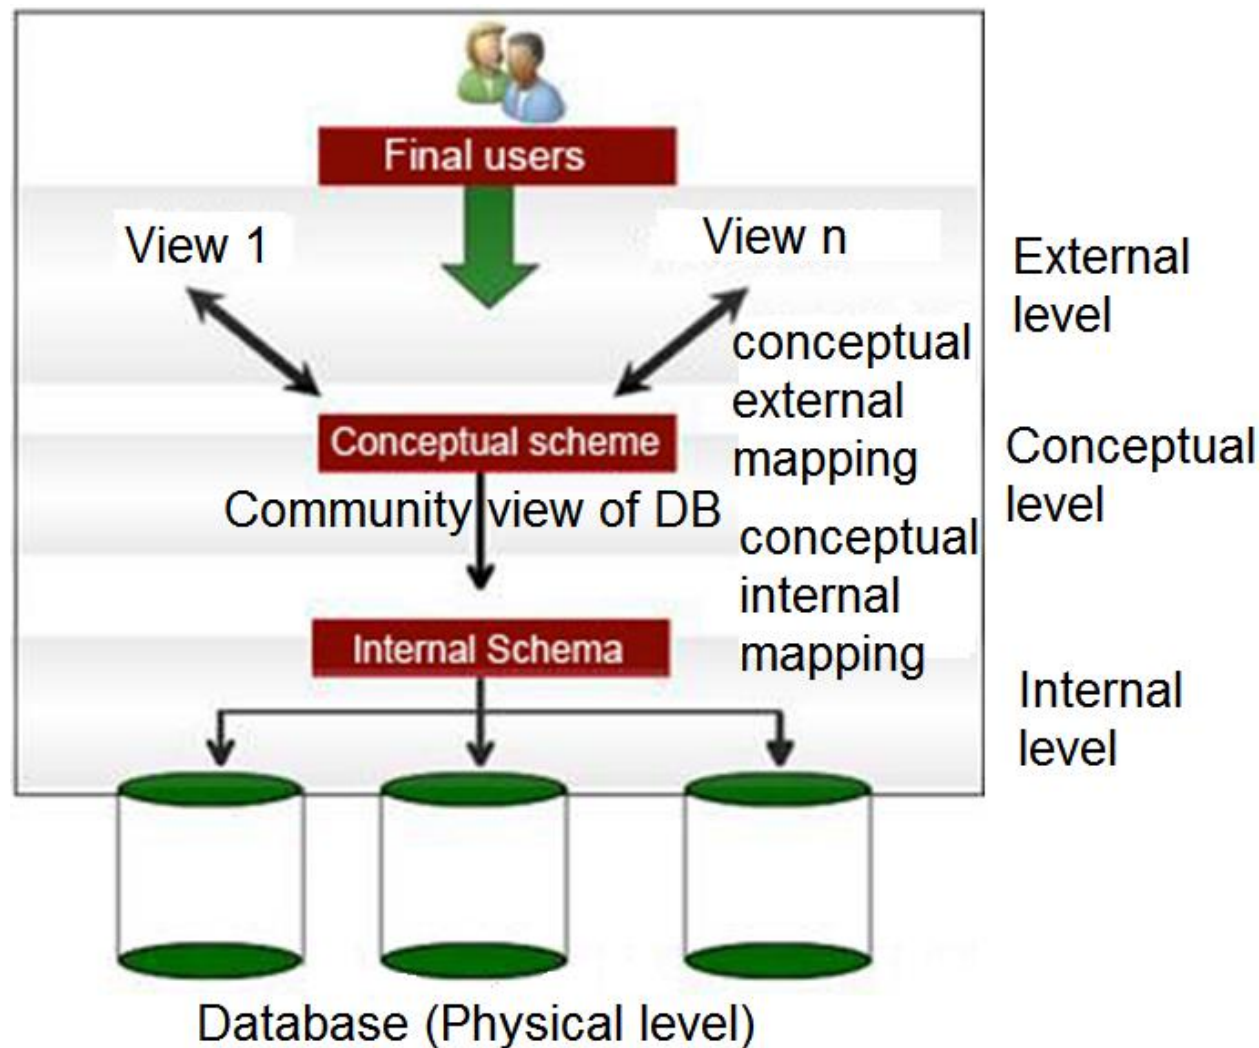
1/ Send a query, something like

```
SELECT *  
FROM Enrollments E INNER JOIN Courses C ON E.CourseTitle = C.CourseTitle  
WHERE StudId = @LoggedInStudentId
```

Issues:

- Mobile developers need to know the database schema + SQL? **YES**
- More details than needed are returned, possibly some sensitive data? **YES**
- What happens if the DA/DBA normalizes the schema, e.g. he/she decides that Departments must have their own table and replaces Department column of Courses with a surrogate FK to Departments table? **NEED TO CHANGE + REPUBLISH THE APP (at least its back-end)**

ANSI/X3 SPARC Architecture for databases



The ANSI/X3 SPARC DBMS Framework: Report of the Study Group on Database Management Systems (1977)

Views

- Physical, conceptual, logical levels
- Why Views?
 - Hide some data from some users
 - Make some queries easier
 - Modularity of database access (customized access = access to parts of the database)
 - Powerful and flexible security mechanism

The bigger the database, more views are used

Views

- A view is a ‘**virtual relation**’ that does not actually exist in the database but is produced upon request, at the time of the request.
- **Base relation** = a named relation corresponding to an entity in the conceptual schema, whose tuples are physically stored in the database.
- **View** = The dynamic result of one or more relational operations operating on the base relations to produce another relation.

Defining and using views

- View $V = \text{Query}(R_1, R_2, \dots, R_n)$ where R_i is a table or another view.
- Schema of V = schema of query result
- “Temporary table”
- In reality, the DBMS re-writes the query Q to use R_1, R_2, \dots, R_n instead of V

Views in SQL

- `CREATE VIEW ViewName AS Query`
- `CREATE VIEW ViewName(A1, A2, ..., An) AS Query`
 - Query - SQL SELECT query
 - Creates the view as an object in the database catalogue; query is run every time the view is opened / used
- `ALTER VIEW` - modifies an existing view
- `DROP VIEW ViewName;`
 - Some DBMS returns an error if View is used in other queries; others returns the error only when the query involving the view is run.
- `SELECT * FROM ViewName` - queries an existing view

Scenario

The mobile dev team of our university develops a student-centred mobile application aiming at helping students to manage their enrollments. How should the app request the list of all enrollments for the logged-in student?

1/ Send a query, something like

```
SELECT *  
FROM Enrollments E INNER JOIN Courses C ON E.CourseTitle = C.CourseTitle  
WHERE StudId = @LoggedInStudentId
```

Issues:

- Mobile developers need to know the database schema + SQL? **YES**
- More details than needed are returned, possibly some sensitive data? **YES**
- What happens if the DA/DBA normalizes the schema, e.g. he/she decides that Departments must have their own table and replaces Department column of Courses with a surrogate FK to Departments table? **NEED TO CHANGE + REPUBLISH THE APP (at least its back-end)**

Scenario

The mobile dev team of our university develops a student-centred mobile application aiming at helping students to manage their enrollments. How should the app request the list of all enrollments for the logged-in student?

2/ Expose data needed by the app through a view

```
SELECT *  
FROM MobileAppView  
WHERE Id = @LoggedInStudentId
```

Issues:

- Mobile developers need to know the database schema + SQL? **NO**
- More details than needed are returned, possibly some sensitive data? **NO**
- What happens if the DA/DBA normalizes the schema, e.g. he/she decides that Departments must have their own table and replaces Department column of Courses with a surrogate FK to Departments table? **THE APP IS NOT TOUCHED**

Scenario

`Courses`(CourseTitle:CHAR(50), Department:CHAR(20), Credits:INTEGER)

`Students`(StudID:INTEGER, StudName:CHAR(50), DoB:DATE, PoB:CHAR(50), Major:CHAR(40), TotalCredits:INTEGER)

`Enrollments`(StudID:INTEGER, CourseTitle:CHAR(50), EnrollmentDate:DATE, Decision:BOOLEAN)

```
CREATE VIEW MobileAppView AS
SELECT StudID AS Id, C.CourseTitle AS Title,
EnrollmentDate, Decision, Department AS DeptName, Credits
FROM Enrollments E
INNER JOIN Courses C ON C.CourseTitle=E.CourseTitle

ALTER VIEW MobileAppView AS
SELECT StudID AS Id, C.CourseTitle AS Title,
EnrollmentDate, Decision, D.Name AS DeptName, Credits
FROM Enrollments E
INNER JOIN Courses C ON E.CourseTitle=E.CourseTitle
INNER JOIN Departments D ON C.DeptID=D.Id
```

Examples

Courses(CourseTitle:CHAR(50), Department:CHAR(20), Credits:INTEGER)

Students(StudID:INTEGER, StudName:CHAR(50), DoB:DATE, PoB:CHAR(50), Major:CHAR(40), TotalCredits:INTEGER)

Enrollments(StudID:INTEGER, CourseTitle:CHAR(50), EnrollmentDate:DATE, Decision:BOOLEAN)

```
CREATE VIEW DBAccepted AS
    SELECT StudID, EnrollmentDate FROM Enrollments
    WHERE CourseTitle='Database' AND Decision='Y'
```

```
CREATE VIEW DBAccepted2 AS
    SELECT Students.StudID, StudName, Major, TotalCredits
    FROM Students, DBAccepted
    WHERE Students.StudID = DBAccepted.StudID
```

Examples

Courses(CourseTitle:CHAR(50), Department:CHAR(20), Credits:INTEGER)

Students(StudID:INTEGER, StudName:CHAR(50), DoB:DATE, PoB:CHAR(50), Major:CHAR(40), TotalCredits:INTEGER)

Enrollments(StudID:INTEGER, CourseTitle:CHAR(50), EnrollmentDate:DATE, Decision:BOOLEAN)

```
CREATE VIEW DBAccepted AS
    SELECT StudID, EnrollmentDate FROM Enrollments
    WHERE CourseTitle='Database' AND Decision='Y'
```

// With Query re-write (v1)

```
CREATE VIEW DBAccepted2 AS
    SELECT Students.StudID, StudName, Major, TotalCredits
    FROM Students,
        (SELECT StudID, EnrollmentDate FROM Enrollments
         WHERE CourseTitle='Database' AND Decision='Y')
        AS DBAccepted
    WHERE Students.StudID = DBAccepted.StudID
```

Examples

Courses(CourseTitle:CHAR(50), Department:CHAR(20), Credits:INTEGER)

Students(StudID:INTEGER, StudName:CHAR(50), DoB:DATE, PoB:CHAR(50), Major:CHAR(40), TotalCredits:INTEGER)

Enrollments(StudID:INTEGER, CourseTitle:CHAR(50), EnrollmentDate:DATE, Decision:BOOLEAN)

```
CREATE VIEW DBAccepted AS
    SELECT StudID, EnrollmentDate FROM Enrollments
    WHERE CourseTitle='Database' AND Decision='Y'
```

// With Query re-write (v2)

```
CREATE VIEW DBAccepted2 AS
    SELECT Students.StudID, StudName, Major, TotalCredits
    FROM Students, Enrollments
    WHERE CourseTitle='Database' AND Decision='Y' AND
        Students.StudID = Enrollments.StudID
```

Examples

Courses(CourseTitle:CHAR(50), Department:CHAR(20), Credits:INTEGER)

Students(StudID:INTEGER, StudName:CHAR(50), DoB:DATE, PoB:CHAR(50), Major:CHAR(40), TotalCredits:INTEGER)

Enrollments(StudID:INTEGER, CourseTitle:CHAR(50), EnrollmentDate:DATE, Decision:BOOLEAN)

```
CREATE VIEW DBAccepted AS
    SELECT StudID, EnrollmentDate FROM Enrollments
    WHERE CourseTitle='Database' AND Decision='Y'
```

// Using views

```
SELECT * FROM DBAccepted WHERE EnrollmentDate < '01-nov-2015'
```

View modification

- Can a view V be modified (insert/delete/update) as any other table?
 - Remember, V is not stored => doesn't make much sense
 - Some users only see the views => it should be possible
- SOLUTION: Modification of V is rewritten (automatically by the system) to modify the base tables.
- Ambiguities: there may be multiple rewrites; which one the one user wanted? Example.

View modification approaches

- Restrict modifications so that the translation to base table modifications is meaningful and unambiguous
 - (+) No user intervention
 - (-) Restrictions may be significant
 - Imposed by SQL standard
- View creator specifies the rewriting process, i.e. what happens in case of delete/update/insert
 - (+) Can handle all modifications
 - (-) No guarantee of correctness
 - Enabled by INSTEAD OF triggers
 - Trigger definition is vendor-specific, not covered by SQL standard

View modification using automatic view modification

- Restrictions in SQL standard for ‘updatable views’:
 - SELECT (no DISTINCT) on a single table T
 - Attributes of T not part of the view should be NULLable or have a default value constraint defined
 - Sub-queries must not refer to T
 - No GROUP BY or HAVING
- Not supported by all DBMS

View modification using automatic view modification

Example: (an updateable view)

```
CREATE VIEW DBAccepted(ID, Curs, EDate) AS
  SELECT StudID, CourseTitle, EnrollmentDate
  FROM Enrollments
  WHERE CourseTitle='Database' AND Decision='Y'
```

- `DELETE FROM DBAccepted WHERE ID=1234`
 - Applies the predicate from view's WHERE AND the predicate from DELETE
- `UPDATE DBAccepted SET EDate = '12-dec-2014' WHERE ID=1234`
- `UPDATE DBAccepted SET Curs = 'Algebra' WHERE ID=1234`
 - Applies the predicate from view's WHERE AND the predicate from UPDATE
- `INSERT INTO DBAccepted (ID, Curs, EDate)`
`VALUES (101, 'Networks', '15-nov-2020')`
 - Record inserted, but the value for Decision column is not set according to the view definition (it is given the default value / NULL)

View modification using automatic view modification

- To avoid insertion of undesired tuples, use WITH CHECK OPTION in view definition and then the previous INSERT is flagged as erroneous
- WITH CHECK OPTION also prevents rows from migrating out of the view, as in UPDATE ... SET Curs = 'Algebra' ...
- WITH [LOCAL/CASCADE] CHECK OPTION - apply /not check on underlying views

View modification using automatic view modification

Example: (an updateable checked view)

```
CREATE VIEW DBAcceptedChecked(ID, Curs, EDate) AS
  SELECT StudID, CourseTitle, EnrollmentDate
  FROM Enrollments
  WHERE CourseTitle='Database' AND Decision='Y'
  WITH CHECK OPTION;
```

Statement	-	WITH CHECK
DELETE FROM DBAccepted WHERE ID=1234	OK	OK
UPDATE DBAccepted SET EDate = '12-dec-2014' WHERE ID=123	OK	OK
UPDATE DBAccepted SET Curs = 'Algebra' WHERE ID=1234	OK	ERROR
INSERT INTO DBAccepted VALUES (101, 'Networks', '2020-01-01')	OK	ERROR

Exercise

Given the following views

```
CREATE VIEW LowSalary AS SELECT * FROM Staff WHERE salary > 9000
```

```
CREATE VIEW HighSalary AS SELECT * FROM LowSalary  
WHERE salary > 10000  
WITH LOCAL CHECK OPTION;
```

```
CREATE VIEW Manager3Staff AS SELECT * FROM HighSalary WHERE  
branch=10;
```

which of the following updates will be rejected/accepted by the DBMS?

- a) UPDATE Manager3Staff SET salary = 9500 WHERE EmpId = 1234
- b) UPDATE Manager3Staff SET salary = 8000 WHERE EmpId = 1234
- c) UPDATE Manager3Staff SET salary = 11000 WHERE EmpId=1234

Exercise

Given the following views

```
CREATE VIEW LowSalary AS SELECT * FROM Staff WHERE salary > 9000
```

```
CREATE VIEW HighSalary AS SELECT * FROM LowSalary  
WHERE salary > 10000  
WITH LOCAL CHECK OPTION;
```

```
CREATE VIEW Manager3Staff AS SELECT * FROM HighSalary WHERE  
branch=10;
```

which of the following updates will be rejected/accepted by the DBMS?

a) UPDATE Manager3Staff SET salary = 9500 WHERE EmpId = 1234
(Rejected)

b) UPDATE Manager3Staff SET salary = 8000 WHERE EmpId = 1234
(Accepted)

c) UPDATE Manager3Staff SET salary = 11000 WHERE EmpId=1234
(Accepted)

Exercise

- Which of the following views are updateable?

a) CREATE VIEW View2 AS

```
SELECT StudName, Major, AVG(TotalCredits)
FROM Students
GROUP BY Major;
```

b) CREATE VIEW View3 AS

```
SELECT DISTINCT Major FROM Students;
```

Exercise

- Which of the following views are updateable?

a) CREATE VIEW View2 AS

```
SELECT StudName, Major, AVG(TotalCredits)
FROM Students
GROUP BY Major;
```

b) CREATE VIEW View3 AS

```
SELECT DISTINCT Major FROM Students;
```

A: None

View modification using triggers

```
CREATE VIEW DBAcceptedUnmodifiable(ID, Name, EDate) AS  
SELECT S.StudID, StudName, EnrollmentDate  
FROM Enrollments E JOIN Students S ON E.StudID = S.StudID  
WHERE CourseTitle='Databases' AND Decision='Y';
```

```
DELETE FROM DBAcceptedUnmodifiable WHERE StudID=1234  
=> error
```

View modification using triggers

```
CREATE TRIGGER DBAcceptedUnmodifiable_OnDelete
INSTEAD OF DELETE ON DBAcceptedUnmodifiable
REFERENCING OLD ROW AS OldRow
FOR EACH ROW
DELETE FROM Enrollments
    WHERE StudID = OldRow.StudID AND
        EnrollmentDate = OldRow.EnrollmentDate AND
        CourseTitle='Database' AND
        Decision='Y'
```

- The WHERE clause in trigger **MUST** match the records to be deleted AND add the expressions of view's WHERE clause (e.g., CourseTitle='Database')
- Tuples of DBAcceptedUnmodifiable don't physically exist, but the Old variable is bind to those that need to be logically deleted!
- The above trigger will only allow deletions; to support UPDATES more trigger(s) are needed!

CAUTION: Writing incorrect trigger will modify/delete unexpected tuples that may even not be part of the view!

Materialized views

- View $V = \text{Query}(R_1, R_2, \dots, R_n)$ where R_i is a table or another view.
- Create a physical table V with schema of query result
- **Execute Query and put results in V**
- Queries refer to V as if it is a table
- (+) Advantage of materialized view: improved query performance
- (-) V can be very large
- (-) Modifications to $R_1, R_2, \dots, R_n \Rightarrow$ recompute/modify V

Advantages of materialized views

- Hide some data from some users
- Make some queries more natural to express
- Modularity of database access
- Improve query performance

Example

- CREATE MATERIALIZED VIEW **MV1**
SELECT StudName, Enrollments.CourseTitle, Credits
FROM Courses, Students, Enrollments
WHERE Enrollments.StudID = Students.StudID AND
Enrollments.CourseTitle = Courses.CourseTitle AND
Students.Major = 'CS'

Note: this is Oracle / PostgreSQL syntax for triggers

Materialized views and modifications

- Modifications to base relations invalidate the view (example)
- DBMS need to maintain the view status
- Modifications on materialized view
 - Just update the stored table
 - Base tables need to be synchronized => same issue with virtual views
- Materialized views are often used to improve performance, hence users will not be allowed to modify them

Design of materialized views

- Which materialized views to create?
- **Efficiency benefits** of materialized views depend on:
 - Size of data
 - Complexity of the view
 - Number of queries using view
 - Number of modifications affecting the view
 - Incremental maintenance vs. Full re-computation
- Analyse the workload using above criteria
- Materialized views generalize the concept of index
- Automatic query rewriting to use materialized views
 - DBMS transparently use existing materialized views to rewrite users' queries, without users even knowing that

Example

```
CREATE MATERIALIZED VIEW CSEnrollments
  SELECT StudID, CourseTitle, Decision FROM Enrollments
  WHERE StudID IN (SELECT StudID FROM Students WHERE
Major='CS')
```

Try to re-write the following query using CSEnrollments :

```
SELECT DISTINCT StudID, TotalCredits
FROM Courses, Students, Enrollments
WHERE Courses.CourseTitle=Enrollments.CourseTitle AND
      Students.StudID=Enrollments.StudID AND
      Courses.Department='INFO' AND Students.Major='CS'
```


Example

```
CREATE MATERIALIZED VIEW CSEnrollments
  SELECT StudID, CourseTitle, Decision FROM Enrollments
  WHERE StudID IN (SELECT StudID FROM Students WHERE
Major='CS')
```

Try to re-write the following query using CSEnrollments :

```
SELECT DISTINCT StudID, TotalCredits
FROM Courses, Students, Enrollments
WHERE Courses.CourseTitle=Enrollments.CourseTitle AND
      Students.StudID=Enrollments.StudID AND
      Courses.Department='INFO' AND Students.Major='CS'
```

```
SELECT DISTINCT StudID, TotalCredits
FROM Courses, CSEnrollments
WHERE Courses.CourseTitle=CSEnrollments.CourseTitle AND
      Courses.Department='INFO'
```

‘Parameterized’ views

- Standard-wise, does not exist
- In general, not available
- MS SQL Server Functions
 - scalar functions: return only scalar/single value; used in SELECT clauses
 - table valued functions (TVF): return a table (set of rows) -> alternative to views (parameterized views)
 - Inline table valued functions (iTVF): contains only one (return) statement that defines the rows/columns to be returned

Inline table valued functions

- Unlike Stored Procedures / multi-statement TVF, the database engine handles this inline TVF as a VIEW
- It computes the execution plan using the statistics on the tables used by this function (similar to views)
- There is no extra load of creating a table variable
- Better in performance than SP / multi-statement TVF

Inline table valued function - Example

-- Definition

```
CREATE FUNCTION EnrolledStudents (@CourseTitle VARCHAR(50))  
RETURNS TABLE  
AS  
RETURN  
SELECT S.*, E.EnrollmentDate, E.Decision  
FROM Students S  
INNER JOIN Enrollments E ON S.StudID = E.StudID  
WHERE E.CourseTitle = @CourseTitle
```

-- Usage

```
SELECT * FROM EnrolledStudents('Database') WHERE Decision='Y'
```


-- Removal

```
DROP FUNCTION EnrolledStudents
```


Bibliography (recommended)

JOAN DESPI
GHEORGHE PETROV

REISZ ROBERT
AUREL STEPAN




Teoria generala a bazelor de date, I. Despi, G. Petrov, R. Reisz, A. Stepan, Mirton, 2000
Cap 11.2.5



A First Course in Database Systems (3rd edition) by Jeffrey Ullman and Jennifer Widom, Prentice Hall, 2007

Chapter 8.1, 8.2, 8.5



Database Systems - A Practical Approach to Design, Implementation, and Management (4th edition) by Thomas Connolly and Carolyn Begg, Addison-Wesley, 2004

Chapter 3.4, 6.4