

Databases 1

Daniel POP

Course Outline

- ~~1. Introduction to database approach~~
- ~~2. The database environment~~
- ~~3. Introduction to The Relational Model~~
- ~~4. Views~~
- ~~5. Transactions~~
- ~~6. SQL Constraints~~
- ~~7. Relational Database Design. Theory and practice~~
- 8. An Introduction to Database Performance. Indexing**
9. JSON Support in Relational Database Management Systems
10. NoSQL Databases

Week 12

Indexing. Basics

Agenda

1. What is an Index
2. Index implementation in relational DBMS
3. Demo

Indexes

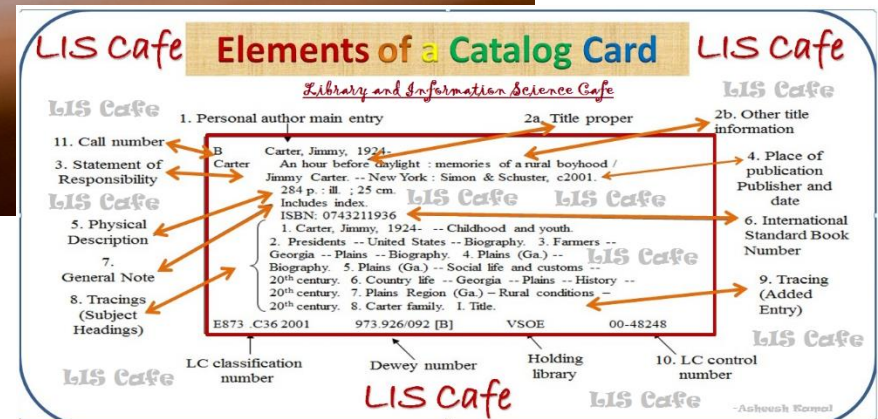
- Primary mechanism to improve the performance of database
- Persistent data structures stored in the database
- Difficult to implement in DBMS
- (Most of the times) Not directly used in queries

Indexes

- Analogy with
 - Index of terms in a book
 - Phonebook
 - Public libraries

SciTech. Ctr.
FA778 .P28
Pattammal, Madras Kanagasabapathy, 1910-
தாய் தன் மகளுக்கு சொல்ல வேண்டிய விஷயங்கள்.
Advice to daughter. எழுதியவர் முட்டை.
பதிப்பு சென்னை, புத்தகம் 1910-
86 p.

CD. DUPL.



Example

Course(CourseTitle:CHAR(50), Department:CHAR(20), Credits:INTEGER)

Student(StudID:INTEGER, StudName:CHAR(50), DoB:DATE, PoB:CHAR(50), Major:CHAR(40))

Enrollment(StudID:INTEGER, CourseTitle:CHAR(50), EnrollmentDate:DATE, Decision:BOOLEAN)

(Name, TotalCredits, ...) SELECT * FROM Student
WHERE StudName='Popescu'

Popescu 50
Ionescu 70
Vasilescu 66
Popescu 45
Ionescu 60

To speed up queries on columns Name (e.g. Student.StudName = 'Popescu') we will build an index on column Name that will quickly return the answer **without scanning the entire table** => Orders of magnitude performance improvements

Example

Course(CourseTitle:CHAR(50), Department:CHAR(20), Credits:INTEGER)

Student(StudID:INTEGER, StudName:CHAR(50), DoB:DATE, PoB:CHAR(50), Major:CHAR(40))

Enrollment(StudID:INTEGER, CourseTitle:CHAR(50), EnrollmentDate:DATE, Decision:BOOLEAN)

| (Name, | TotalCredits, ...) | SELECT StudID, StudName FROM Student WHERE StudName = 'Popescu' AND TotalCredits > 60 |
|-----------|--------------------|--|
| ----- | | |
| Popescu | 50 | |
| Ionescu | 70 | |
| Vasilescu | 66 | |
| Popescu | 45 | |
| Ionescu | 60 | |

Example

Course(CourseTitle:CHAR(50), Department:CHAR(20), Credits:INTEGER)

Student(StudID:INTEGER, StudName:CHAR(50), DoB:DATE, PoB:CHAR(50), Major:CHAR(40))

Enrollment(StudID:INTEGER, CourseTitle:CHAR(50), EnrollmentDate:DATE, Decision:BOOLEAN)

| (Name, | TotalCredits, ...) | |
|-----------|--------------------|---|
| ----- | | |
| Popescu | 50 | SELECT StudID, StudName FROM Student WHERE StudName = 'Popescu' AND TotalCredits > 60 |
| Ionescu | 70 | |
| Vasilescu | 66 | SELECT StudName, CourseTitle FROM Student JOIN Enrollment ON Student.StudID = Enrollment.StudID |
| Popescu | 45 | |
| Ionescu | 60 | |

Index Implementation

- Balanced B-Tree
- Each node is represented by a 8k page
- Pages on same level -> double linked list
- Logarithmic running time
- Very good for:
 - Searching for a value “=”
 - Scanning a range of values “BETWEEN” ($\text{value1} < A < \text{value2}$)
 - Sorted output (table joins)

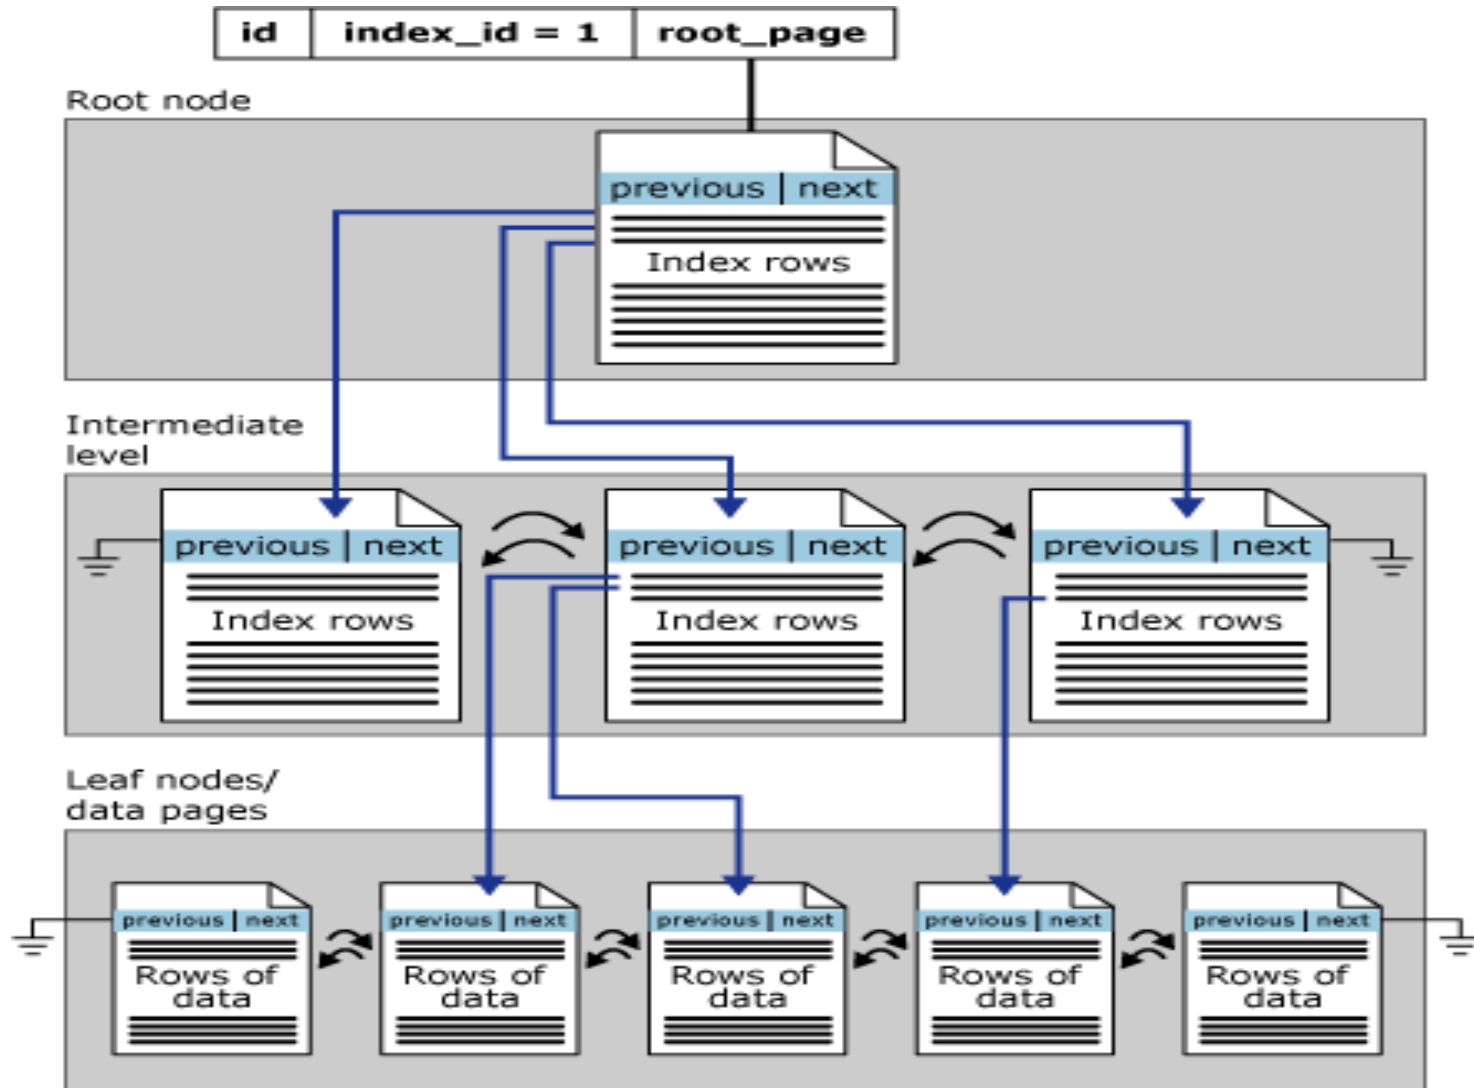
Reminder

A balanced B-Tree requires that each leaf is at the same distance from the root.

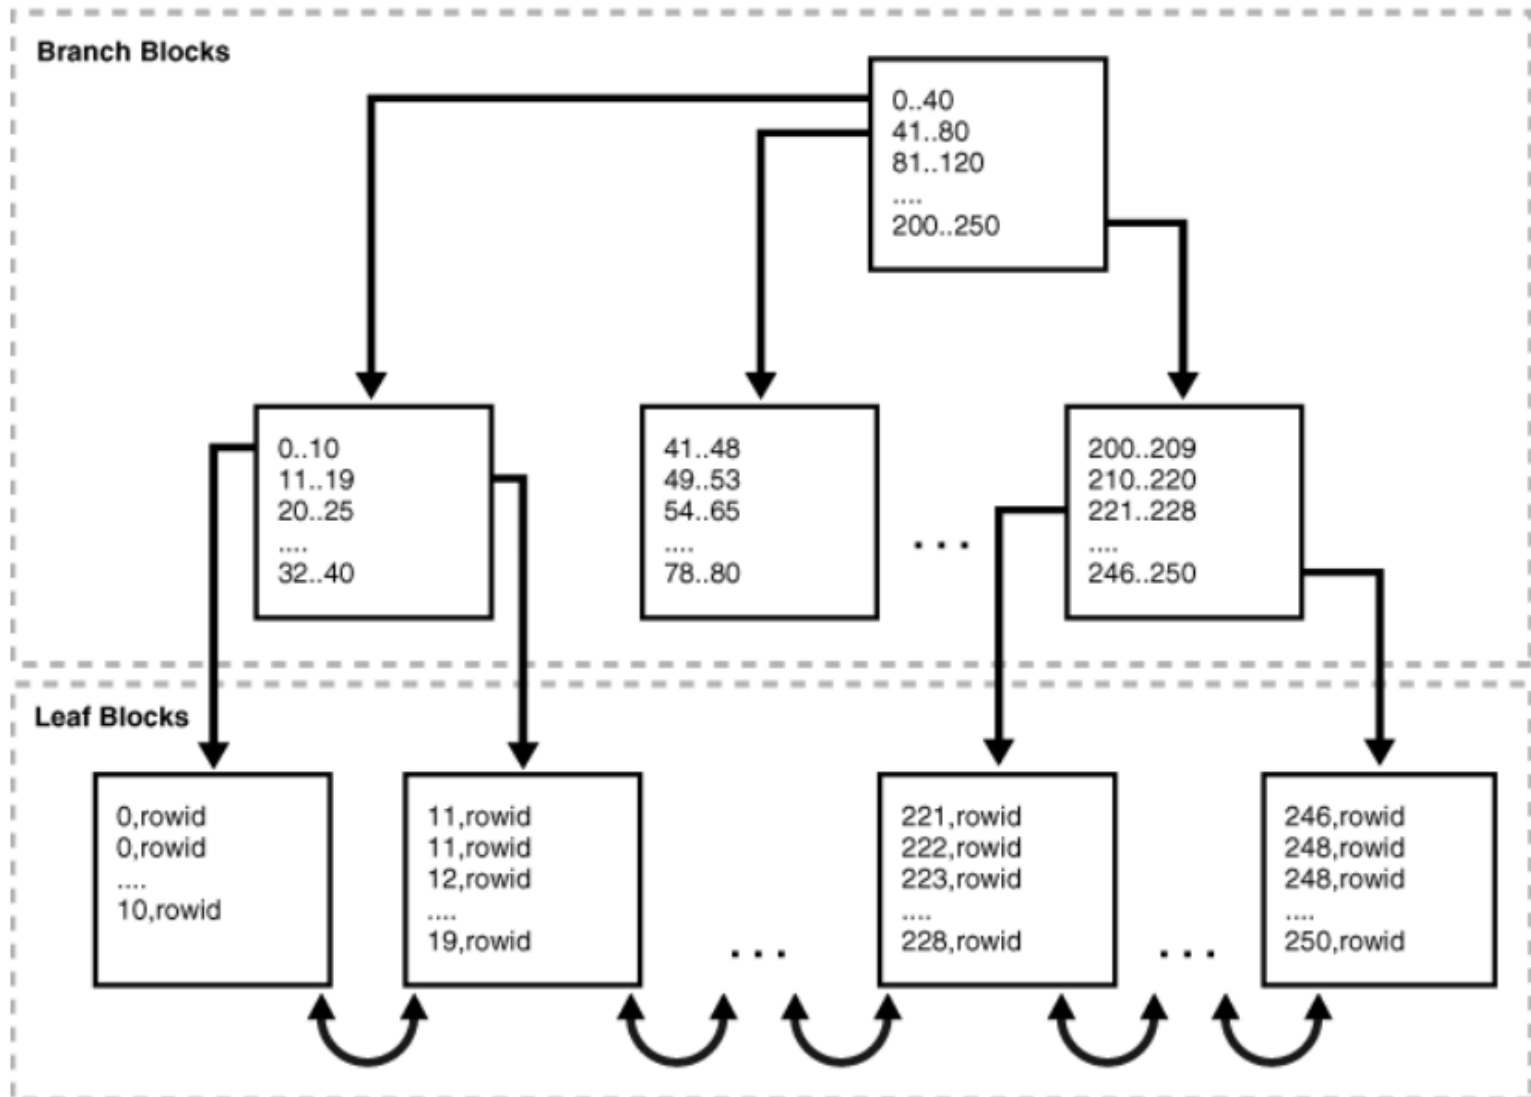
MS SQL Server Clustered vs. Non-clustered

- A clustered index
 - Type of index that reorders the way records in the table are physically stored
 - Therefore table can have only one clustered index.
 - The leaf nodes of a clustered index contain the data pages.
- A non-clustered index
 - Type of index in which the logical order of the index does not match the physical stored order of the rows on disk.
 - A table may have multiple non-clustered indexes
 - The leaf node of a non-clustered index does not consist of the data pages. Instead, the leaf nodes contain index rows.

MS SQL Server Clustered Index



MS SQL Server Non-Clustered Index



Description of "Figure 3-1 Internal Structure of a B-tree Index"

SQL Server Clustered vs. Non-clustered

- A clustered index means you are telling the database to store close values actually close to one another on the disk. This has the benefit of rapid scan / retrieval of records falling into some range of clustered index values.
- Example: If you wish to quickly retrieve all courses on which a student is enrolled to, you may wish to create a clustered index on *StudID* column of the Enrollment table. This way the records with the same StudID will be physically stored close to each other on disk (clustered) which speeds up their retrieval.
- Remark: The index on StudID will obviously be not unique, so you either need to add a second field to "uniquify" the index or let the database handle that for you.

SQL Server Clustered vs. Non-clustered

- The key difference between clustered indexes and non clustered indexes is that the leaf level of the clustered index is the table. This has two implications
 - The rows on the clustered index leaf pages always contains something for each of the (non sparse) columns in the table (either the value, or a pointer to the actual value).
 - The clustered index is the primary copy of a table.
- Non clustered indexes can also store actual data by using the INCLUDE clause (Since SQL Server 2005) to explicitly include non key columns
- If a table has no clustered index it is called a heap. Non-clustered indexes can be created on both heap and clustered tables.
- When you create a PRIMARY KEY constraint, a unique clustered index on PK column(s) is automatically created if a clustered index on the table does not already exist.

<https://stackoverflow.com/questions/1251636/what-do-clustered-and-non-clustered-index-actually-mean>

Clustered vs. Non-clustered - Other DBMS

SQL SERVER

- SQL Server uses clustered indexes by default (index-organized tables), using the primary key as clustering key.
- To create a heap table you must use the NONCLUSTERED clause in the primary key definition

ORACLE

- Oracle database uses heap tables by default. Index-organized tables can be created using the ORGANIZATION INDEX clause, always using the primary key as the clustering key.

PostgreSQL

- PostgreSQL only uses heap tables
- You can, however, use the CLUSTER clause to align the contents of the heap table with an index.

MySQL

- The MyISAM engine only uses heap tables while the InnoDB engine always uses clustered indexes.


<https://use-the-index-luke.com/sql/clustering/index-organized-clustered-index>

MS SQL Server - Hash Index

- Only on Memory-Optimized tables
- Fixed size
- Deterministic
- Poisson or bell curve distribution
- Optimized for point lookups “=”
- Constant running time
- Not good for range scans or inequality clauses
- If on avg > 100 entries/value then use non-clustered index
- Number of buckets must be set when index is created

<https://docs.microsoft.com/en-us/sql/database-engine/determining-the-correct-bucket-count-for-hash-indexes?view=sql-server-2014>


Hash table



Reminder
Hash function ensures a uniform distribution of hash values.

Hash collision strategy: in this example, separate chaining with head records in the bucket array.

MS SQL Server - Columnstore Index



- Can only be scanned
- High compression rates (up to 10x)
- No sort order
- Segment and column elimination
- Good for:
 - Scanning
 - Aggregation

<https://www.red-gate.com/simple-talk/databases/sql-server/t-sql-programming-sql-server/what-are-columnstore-indexes>

<https://learn.microsoft.com/en-us/sql/relational-databases/indexes/columnstore-indexes-overview?view=sql-server-ver16>

<https://swarm64.com/post/postgresql-columnstore-index-intro/>

Exercise

- Given the following query

```
SELECT StudID, CourseTitle, Credits  
FROM Enrollment, Course  
WHERE Enrollment.CourseTitle = Course.CourseTitle AND  
Enrollment.Decision = TRUE and Course.Credits > 4
```

which of the following indexes could NOT be useful in speeding-up the query?

- a) tree-based index on Enrollment.CourseTitle
- b) hash-based index on Enrollment.Decision
- c) hash-based index on Course.CourseTitle
- d) hash-based index on Course.Credits

Shortcomings of indexes

- Extra space on database
- Overhead of index creation (usually when the database is loaded)
- Maintenance of index: index has to be modified every time values in the table are changed so that it slows down the modifications (add/delete/update)
 - Defragment
 - Rebuild

SQL Implementation

- Indexes are created by default for PRIMARY KEY and UNIQUE constraints
- CREATE INDEX IdxName ON Student (Name)
- CREATE UNIQUE INDEX IdxCNP ON Student (CNP)
 - Check that all values for CNP are unique and will generate an error in case of duplicates
- ALTER INDEX IdxName ON TableName REBUILD
- DROP INDEX IdxName

MS SQL Server Implementation

```
CREATE CLUSTERED INDEX [CLI_SalesOrderDetailID]  
ON [dbo].[BigTable] ([SalesOrderDetailID] ASC)
```

```
CREATE NONCLUSTERED INDEX [NonClusteredIndexDemo]  
ON [dbo].[BigTable]  
(  
    [UnitPrice] ASC,  
    [ModifiedDate] ASC  
)  
INCLUDE ([OrderQty])
```


Demo

- Adventure Works Sample Database from MS
- BigTable example
 - Structure. Number of records
 - Indexes
- Performance demo
 - CPU time vs. Elapsed time
 - Discuss the query plan
 - Look into physical structure of the index

Demo

```
-- Enable the execution time tracking in the messages tab  
SET STATISTICS TIME ON
```

```
-- Index Seek vs. Table Scan (clustered index) Performance
```

```
SELECT OrderQty FROM [dbo].[BigTable]  
WITH (INDEX(CLI_SalesOrderDetailID))  
WHERE UnitPrice = 5.70
```

```
SELECT OrderQty FROM [dbo].[BigTable]  
WITH (INDEX(NonClusteredIndexDemo))  
WHERE UnitPrice = 5.70
```

```
-- Aggregating a column using a column-store index vs. clustered index
```


```
SELECT AVG(UnitPrice) FROM [dbo].[BigTable]  
WITH (INDEX(CLI_SalesOrderDetailID))
```

```
SELECT AVG(UnitPrice) FROM [dbo].[BigTable]  
WITH (INDEX(NonClusteredColumnStoreIndexDemo))
```


Bibliography (recommended)

JOAN DESPI
GHEORGHE PETROV


REISZ ROBERT
AUREL STEPAN



Teoria generala a bazelor de date,
I. Despi, G.
Petrov, R. Reisz,
A. Stepan,
Mirton, 2000
Cap 10.4 - 10.5



A First Course in Database Systems (3rd edition) by Jeffrey Ullman and Jennifer Widom, Prentice Hall, 2007
Chapter 7.1, 7.2, 7.3 – SQL Constraints
Chapter 8.3 – 8.4 - Indexes



Database Systems - A Practical Approach to Design, Implementation, and Management (4th edition) by Thomas Connolly and Carolyn Begg, Addison-Wesley, 2004
Chapter 17.3 & C5